

Appendix C:

Automatic Categorization of Questions for a Mathematics Education Service

Ken Williams Rafael A. Calvo

David Bell

Web Engineering Group

The University of Sydney

Bldg J03, Sydney NSW 2006

{kenw,rafa}@ee.usyd.edu.au, dave@student.usyd.edu.au

Abstract. This paper describes a new approach to managing a stream of questions about mathematics by integrating a text categorization framework into a relational database management system. The corpus studied is based on unstructured submissions to an ask-an-expert service in learning mathematics. The classification system has been tested using a Naïve Bayes learner built into the framework. The performance results of the classifier are also discussed. The framework was integrated into a PostgreSQL database through the use of procedural trigger functions.

1 Introduction

Ask-an-expert services are becoming more common, spanning from standard customer relationship management to discussion forums in a particular discipline. In general, these online services are supported by domain experts who attempt to answer questions posted via email or web forms. Since these experts often have a single subdomain of expertise it is very helpful if they have only to read questions that relate to this subdomain. This can be done by organizing the service in such a way that users are encouraged to post their question in the appropriate area. However, this approach is not always successful as often the user will either ignore the organization scheme or not know to which area their question belongs.

These problems are common within a number of domains. Our test was performed on messages sent to a mathematics ask-an-expert service for students and teachers.[6] The issues discussed also apply to other similar systems such as customer relationship management (CRM) and e-learning systems in general. These systems can use an automatic text categorization framework to categorize the questions into the experts' areas of interest, or into the appropriate customer support mailbox.

The downside of an automatic categorization approach is that integrating such functionality into existing systems can be very complex, and often involves an in depth understanding of text categorization techniques. Also, the content is normally stored in systems with a relational database in the backend, as is the case for most content and learning management systems. By building the categorizer into the database, the categorization framework[11] can be made invisible to the users and is thus more attractive to the average system administrator

Implicit Differentiation

Find the slope of the tangent at the point
(3,4) on the circle $x^2 + y^2 = 25$.

My answer: I guess we would need to put it
in the $y = mx + b$ form.

Thanks for any help,
Scott

Figure 1: An example document from the Dr. Math corpus

or application developer. Also, application developers, do not have to re-implement the classification software. They only need a machine learning professional to assist in training the classifier, and once trained it can then be reused for different applications.

The applications of information retrieval have been well studied since the 1980s, as discussed by Salton [9, 8], and many of these methodologies have been integrated into commercial database management systems that have free text search capabilities. However, this integration does not seem to have penetrated the text categorization domain.

Section 2 of the paper discusses the data set that was used to test the system. Section 3 discusses the text categorization framework and the extensions made to it, including the implementation within the database management system. Section 4 discusses the quantitative results of the testing process and Section 5 concludes.

2 Dr. Math Corpus

For the evaluation of our system we have tested the performance of the categorization system over a set of unstructured, informal documents from the Ask Dr. Math service.[6] These documents are mostly written by students between the ages of 6 and 18, though question submissions can come from any member of the general public. The documents vary in length from a single sentence to several paragraphs. In addition to this, many examples contain symbols and diagrams, making linguistic analysis very difficult. The Ask Dr. Math service has about 300 volunteers (about 30-40 of which may be active in any given month), dealing with hundreds of questions a day. The volunteers have expertise in different areas of math, and the site has won a number of awards for its useful service. Figure 1 shows an example submission to the service.

The filtering of questions is a major element of the Ask Dr. Math question answering process. The service may receive about 7000 questions a month, about half of which are eventually answered. The unanswered questions may be duplicate submissions, messages of thanks, inappropriate questions, or other messages that don't require a response. There also may be some legitimate mathematics questions that go unanswered, simply because the service is not fee-based for either the students or the experts, and thus can make no guarantee that any particular question will be answered. The experts are currently responsible for choosing their own questions to answer.

The Dr. Math corpus we used contains 6632 documents and was split into a training




Figure 2: Category distribution for the 95 Dr. Math categories

set of 5305 documents and a testing set of 1327 documents. There are 95 categories in the corpus, and the average number of documents in each category is 107.15. The categories are structured in a 2-level hierarchy (“level” and “topic”), with a total of four levels (elementary school, middle school, high school, and college) and 62 topics (for example, “geometry” or “statistics”). The most popular category, high school-level geometry, contains 877 documents, and the least popular category, elementary-level golden ratio, contains only 3 documents. Each document may be a member of more than one category, and the average number of categories per document is 1.53. Figure 2 shows the distribution of categories throughout the corpus.

It may be important to note that our corpus was drawn from the public archive of answered questions, not directly from the stream of incoming questions. Because the archiving process is fairly intensive and not all questions are chosen for archiving, our corpus may therefore differ significantly from the incoming question stream. For example, none of the kinds of unanswered questions mentioned earlier are represented in the archive. Because of this difference, it is difficult to extrapolate our experiments to performance on the incoming question stream. However, because the incoming question stream is uncategorized, obtaining a large enough number of categorized questions for our investigation necessitated drawing them from the archived questions.

3 Categorization Framework

Object Oriented Application Frameworks (OOAF) are software engineering artifacts that improve reusability of design and implementation [4, 5].

The framework used in this project was designed to provide a consistent way to build document categorization systems.[11] It allows the developer to focus on the quality of the more critical and complex modules by allowing reuse of common, simple base modules. The

```
INSERT INTO documents
  (name, content, categories)
VALUES
  ('my name',
   'my content',
   categorize('my name',
              'my content',
              'documents'));
```

Figure 3: Example document insertion statement with categorization

framework has implementations of k-Nearest-Neighbor (kNN), Naïve Bayes (NB), Support Vector Machine (SVM), and Decision Tree (DT) classifiers [12, 10]. Other methods such as Neural Network [3, 2] classifiers are under development.

The framework architecture allows extensions to be built by subclassing its main classes [11]. Class inheritance contributes to code reuse and quality. In this project we extended the framework by adding an alternative Collection class to allow for the data to be read directly from a database instead of from a file system. Having a classifier that uses the data directly from the database streamlines the management of questions and answers in this type of system. In fact, it allows many content or learning management systems to natively use automatic classification features. The framework also provides statistical analysis of experimental results, and produced the performance measures discussed in Section 4.

The framework's architecture and language choice enabled us to easily build the framework into PostgreSQL through PostgreSQL's PL/Perl and PL/perlU support. This support allows the creation of procedural language functions through the use of the "CREATE FUNCTION" SQL command. Using this support and the PL/perlU language we were able to build a "launching" function that invoked the categorization framework on the document to be classified. This means that the only command necessary to categorize a document is a basic insert statement with a function call in place of a value for the category of the document, as shown in Figure 3. This statement can be further simplified through the creation of a pl/sql trigger function which fires automatically on insertion and passes the necessary values to the `categorize()` function.

If the categorization is to take place within a database, where categorized documents are often going to be appended to the learning set, a learning algorithm which has very little training overhead is ideal. This avoids the need to retrain a categorizer after each document insertion. Two such algorithms are the Naïve Bayes algorithm (NB) and the K-Nearest Neighbor algorithm (kNN).

The training phase in NB consists only of counting term frequencies in each document and using them to calculate conditional probabilities between terms and categories. These probabilities are then consulted when categorizing a new document, with conditional probabilities for each term being multiplied to find the probability that a given document belongs to a certain category.

kNN in its basic form has essentially no training phase. Each document is represented as an n -dimensional vector, where n is the number of unique terms in the training set. When a new document is to be classified, it is compared to the vectors of the documents in the training set. The k training vectors which are closest to the test vector are found (with distance

	MaP	MaR	MaF_1	MiF_1	Error
NB	0.246	0.226	0.226	0.361	0.022
kNN	0.211	0.186	0.179	0.257	0.025
Baseline	0.019	0.018	0.018	0.042	0.031

Table 1: Macro- and Micro-averaged performance scores.

defined as the cosine of the angle between any two vectors), and the most prevalent category or categories amongst these is assigned to the new document.

In our testing, we have found that NB is a more accurate categorizer than kNN. The rest of this paper will focus on the NB experiment and results.

Since the categorization performance is determined only by the classification framework, all these methods should behave the same inside the database as they do outside the database. What the integration into the database does is to make the functionalities of the framework available as procedures in the SQL language. Since relational databases can be designed using an object oriented methodology [1, 7], by integrating it in this way, the classification task (and framework) can also be designed into larger OO systems.

4 Method and Results

The 5305 training documents in the Dr. Math corpus were loaded into a database table named “documents.” This table consisted of 3 columns: name, content and categories. The testing set was then inserted into the database, one document at a time using a statement similar to that in Figure 3. A SELECT statement was then used to compare the assigned and actual categories of each document. Through this comparison the performance of the categorization in terms of precision and recall was measured.¹ The precision and recall were then used to calculate the F_1 measure [2, 10].

The results in Table 1 show the performance of the categorization algorithms. The precision, recall, and F_1 scores can be computed using macro-averaging, which gives equal weight to each category, or micro-averaging, which gives equal weight to each document. [12, 10] For the kNN algorithm, we used a k value of 15 and a categorization threshold of 0.12, which seemed to perform the best in our investigations. We also include in Table 1 the results of a baseline classifier, which assigns categories at random to each test document, weighting the random generator by the frequency of categories in the training set.

Next, we turned our attention to ways of improving performance on the test set. As mentioned in Section 2, each category name is a combination of two components, a level and a topic. This suggests that separate categorizers could be trained to recognize the two components separately, perhaps with more success than a single categorizer may have on the two components together.

To test this hypothesis, we created three new categorization tasks: one that categorizes by level alone, one that categorizes by topic alone, and one that uses the separate topic and level categorizers to assign a combined category. In this combined process, each assigned level was combined with each assigned topic, any nonexistent categories (such as “calculus.elem”

¹Recall is the proportion of the target items that the system selected, i.e. $tp/(tp+fn)$. Precision is the proportion of selected items the system got right, i.e. $tp/(tp+fp)$.

Task	MaP	MaR	MaF_1	MiF_1	Error
Level	0.524	0.626	0.570	0.671	0.223
Topic	0.339	0.314	0.313	0.440	0.026
Both	0.187	0.181	0.166	0.223	0.035

Table 2: Performance of Naïve Bayes classifier on subtasks.

Task	MaP	MaR	MaF_1	MiF_1	Error
Level	0.326	0.319	0.322	0.468	0.328
Topic	0.029	0.027	0.027	0.067	0.041
Both	0.015	0.010	0.011	0.035	0.027

Table 3: Performance of baseline classifier on subtasks.

or “addition.college”) were filtered out, and all remaining categories were assigned to the given document. This process was performed using the Naïve Bayes and baseline categorizers described above. Table 2 shows the performance for the Naïve Bayes categorizer and Table 3 shows the baseline categorizer for comparison.

Comparing the combined task to the NB results in Table 1, we see that separating the categorization task into two subtasks adversely affected the overall performance on the combined task. The performance for the level task seems good, but comparing it with the baseline categorizer shows that it may not be significantly better than random guessing. This is probably due to the small number of categories. However, the performance on the topic task is noteworthy, because it is so far above both the baseline categorizer and the original NB categorizer. In addition, the topic assignment may be more valuable in this application than the level assignment, because while most students will be able to indicate their own age or grade level when asking a question, they may not be able to place their own question in an appropriate category.

Because the math topics used in this experiment were generated from the category names, we ended up with some duplication that may have decreased performance. For example, there are categories called “probability.high,” “statistics.high,” and “prob.stats.middle.” This means that the combined list of extracted math topics includes “probability,” “statistics,” and “prob.stats.” The exact effect of this on the categorizers is unknown, but we might expect these overlapping categories to confuse the categorizer. A service such as this one may therefore benefit from using more consistent category names.

5 Conclusion

We have described a system that integrates a categorization framework into a relational database. The results show it is possible to integrate categorization techniques into the relational databases used by learning and content management systems.

Two categorization algorithms were applied to the task of classifying messages sent to an educational ask-an-expert service. A Naïve Bayes classifier outperformed the kNearest Neighbour classifier and was reasonably successful at categorizing the messages. Future work includes testing classifiers that use other machine learning models such as Support Vector Machines and Neural Networks. The classification performance of Naïve Bayes was also

measured using the 2-level hierarchy of the corpus. In this case, the highest success rate was found in classifying messages by math topic, instead of school level. This task could be useful on the unstructured data sent to ask-an-expert services such as the one we discussed.

Acknowledgements

D. Bell gratefully acknowledges support from the Capital Markets Collaborative Research Centre and the University of Sydney. The authors also acknowledge The Math Forum for giving permission to use and discuss their data.

References

- [1] Michael R. Blaha, William J. Premierlani, and James E. Rumbaugh. Relational database design using an object-oriented methodology. *Communications of the ACM*, 31(4):414–427, 1988.
- [2] Rafael A. Calvo. Classifying financial news with neural networks. In *6th Australasian Document Symposium*, page 6, December 2001.
- [3] Rafael A. Calvo and H. A. Ceccatto. Intelligent document classification. *Intelligent Data Analysis*, 4(5), 2000.
- [4] Mohamed Fayad and Douglas C. Schmidt. Object-oriented application frameworks. *Communications of the ACM*, 40(10):32–38, 1997.
- [5] Mohamed Fayad and Douglas C. Schmidt, editors. *Building Application Frameworks*. John Wiley & Sons, 1999.
- [6] The Math Forum. Ask Dr. Math web site. <http://www.mathforum.org/dr.math/>.
- [7] James Rumbaugh, Michael Blaha, William Premierlani, Frederick Eddy, and William Lorensen. *Object-oriented modeling and design*. Prentice-Hall, Inc., 1991.
- [8] Gerald Salton. Developments in automatic text retrieval. *Science*, (253):974–979, 1991.
- [9] Gerard Salton. *Automatic Text Processing: The Transformation, Analysis, and Retrieval of Information by Computer*. Addison-Wesley, Reading, Pennsylvania, 1989.
- [10] Fabrizio Sebastiani. Machine learning in automated text categorization. *ACM Computing Surveys (CSUR)*, 34(1):1–47, 2002.
- [11] Ken Williams and Rafael A. Calvo. A framework for text categorization. *7th Australasian Document Computing Symposium*, 2002.
- [12] Yiming Yang and X. Liu. A re-examination of text categorization methods. In *22nd Annual International SIGIR*, pages 42–49, Berkley, August 1999.